DocChat AI Architecture Overview - Page 1
DocChat AI is an enterprise-grade document intelligence platform powered by Retrieval-Augmented Gene
The DocumentLoader service is responsible for ingesting files from in-memory byte buffers or the local file

DocChat AI Chunking & Embedding Pipeline - Page 2

Text chunking is executed by the TextChunker service. Rather than arbitrarily slicing strings at rigid byte b
This overlap prevents semantic cliff effects where critical context or named entities are split across retrieva
Each generated chunk is enriched with essential metadata: doc_id, filename, source page number, global

DocChat AI Vector Storage & Persistence - Page 3

Once vector embeddings are generated, chunks and their metadata are committed to a persistent Chroma

The ChromaVectorStore exposes a clean public interface featuring add_chunks() for upserting vectorized

This ingestion architecture guarantees high retrieval accuracy, sub-50ms vector query latencies, and deter